

```
!pip install keras keras-hub --upgrade -q
```

```
0.0/1.4 MB ? eta -:-:--  
1.4/1.4 MB 43.2 MB/s eta  
0:00:01 1.4/1.4 MB 28.9 MB/s  
eta 0:00:00
```

```
0.0/947.9 kB ? eta -:-:--  
947.9/947.9 kB 56.4 MB/s eta
```

```
0:00:00
```

ERROR: pip's dependency resolver does not currently take into account all the packages that are installed. This behaviour is the source of the following dependency conflicts.

keras-nlp 0.21.1 requires keras-hub==0.21.1, but you have keras-hub 0.22.2 which is incompatible.

```
import os  
os.environ["KERAS_BACKEND"] = "jax" # Or "tensorflow" or "torch"
```

```
import keras  
from keras import layers  
import keras_hub  
from keras.utils import image_dataset_from_directory  
import numpy as np  
import matplotlib.pyplot as plt  
import pathlib  
import shutil  
import zipfile  
import pandas as pd  
import random  
import tensorflow as tf
```

```
zip_file_path = "/content/cats_vs_dogs_small.zip"
```

```
if os.path.exists(zip_file_path):  
    print(f"Found {zip_file_path}, now extracting...")  
    with zipfile.ZipFile(zip_file_path, "r") as zip_ref:  
        zip_ref.extractall("/content/")  
    print("Data extracted successfully.")  
else:  
    print(f" '{zip_file_path}' not found. Please upload the zipped  
folder first.")  
    raise FileNotFoundError(f"'{zip_file_path}' not found.")
```

```
# Define the base directory path after unzipping
```

```
base_dir = pathlib.Path("/content/cats_vs_dogs_small")  
source_train_dir = base_dir / "train"  
source_val_dir = base_dir / "validation"  
source_test_dir = base_dir / "test"
```

```

if base_dir.exists():
    num_train = len(list(source_train_dir.glob("*/*.jpg")))
    num_val = len(list(source_val_dir.glob("*/*.jpg")))
    num_test = len(list(source_test_dir.glob("*/*.jpg")))
    print(f"\nData verification successful!")
    print(f"  Total original training images found:  {num_train}")
    print(f"  Total original validation images found: {num_val}")
    print(f"  Total original test images found:       {num_test}")
else:
    print(f" Base directory not found at {base_dir}")
    raise FileNotFoundError("Base directory not found after
unzipping.")

```

Found /content/cats_vs_dogs_small.zip, now extracting...
Data extracted successfully.

Data verification successful!
Total original training images found: 2000
Total original validation images found: 1000
Total original test images found: 1000

```

def create_experimental_training_set(source_dir, dest_dir,
train_size_per_class):

    if dest_dir.exists():
        shutil.rmtree(dest_dir)
        os.makedirs(dest_dir)

    for category in ("cats", "dogs"):
        os.makedirs(dest_dir / category, exist_ok=True)
        fnames = [f for f in os.listdir(source_dir / category) if
f.endswith('.jpg')]
        random.shuffle(fnames)

        num_to_copy = min(train_size_per_class, len(fnames))
        for i in range(num_to_copy):
            src = source_dir / category / fnames[i]
            dst = dest_dir / category / fnames[i]
            shutil.copyfile(src, dst)
        print(f"Created experimental training set with {num_to_copy}
images per class.")

def create_experimental_validation_set(source_dir, dest_dir,
val_size_per_class):

    if dest_dir.exists():
        shutil.rmtree(dest_dir)
        os.makedirs(dest_dir)

    for category in ("cats", "dogs"):

```

```

        os.makedirs(dest_dir / category, exist_ok=True)
        fnames = [f for f in os.listdir(source_dir / category) if
f.endswith('.jpg')]
        random.shuffle(fnames)

        num_to_copy = min(val_size_per_class, len(fnames))
        for i in range(num_to_copy):
            src = source_dir / category / fnames[i]
            dst = dest_dir / category / fnames[i]
            shutil.copyfile(src, dst)
        print(f"Created experimental validation set with {num_to_copy}
images per class.")

def create_experimental_test_set(source_dir, dest_dir,
test_size_per_class):

    if dest_dir.exists():
        shutil.rmtree(dest_dir)
    os.makedirs(dest_dir)

    for category in ("cats", "dogs"):
        os.makedirs(dest_dir / category, exist_ok=True)
        fnames = [f for f in os.listdir(source_dir / category) if
f.endswith('.jpg')]
        random.shuffle(fnames)

        num_to_copy = min(test_size_per_class, len(fnames))
        for i in range(num_to_copy):
            src = source_dir / category / fnames[i]
            dst = dest_dir / category / fnames[i]
            shutil.copyfile(src, dst)
        print(f"Created experimental test set with {num_to_copy} images
per class.")

data_augmentation_layer = keras.Sequential(
    [
        layers.RandomFlip("horizontal"),
        layers.RandomRotation(0.1),
        layers.RandomZoom(0.2),
    ],
    name="data_augmentation"
)

def build_scratch_model(input_shape=(180, 180, 3), dropout_rate=0.25):
    inputs = keras.Input(shape=input_shape)
    # Augmentation as a layer
    x = data_augmentation_layer(inputs)
    # Rescaling
    x = layers.Rescaling(1.0 / 255)(x)

```

```

    # Rest of the model
    x = layers.Conv2D(32, 3, activation="relu")(x); x =
layers.MaxPooling2D(2)(x)
    x = layers.Conv2D(64, 3, activation="relu")(x); x =
layers.MaxPooling2D(2)(x)
    x = layers.Conv2D(128, 3, activation="relu")(x); x =
layers.MaxPooling2D(2)(x)
    x = layers.Conv2D(256, 3, activation="relu")(x); x =
layers.MaxPooling2D(2)(x)
    x = layers.Conv2D(512, 3, activation="relu")(x); x =
layers.GlobalAveragePooling2D()(x)
    x = layers.Dropout(dropout_rate)(x)
    outputs = layers.Dense(1, activation="sigmoid")(x)
    model = keras.Model(inputs=inputs, outputs=outputs)
    model.compile(loss="binary_crossentropy", optimizer="adam",
metrics=["accuracy"])
    return model

def build_pretrained_xception_model(trainable_base=False,
dropout_rate=0.25, input_shape=(180, 180, 3)):

    try:

        conv_base =
keras_hub.models.Backbone.from_preset("xception_41_imagenet",
trainable=trainable_base, name="xception_backbone")
        preprocessor =
keras_hub.layers.ImageConverter.from_preset("xception_41_imagenet",
image_size=(180, 180))

        inputs = keras.Input(shape=input_shape)
        # Apply augmentation first
        x = data_augmentation_layer(inputs)

        x = preprocessor(x)

        x = conv_base(x)
        # Add classification head
        x = layers.GlobalAveragePooling2D()(x)
        x = layers.Dense(256, activation="relu")(x)
        x = layers.Dropout(dropout_rate)(x)
        outputs = layers.Dense(1, activation="sigmoid")(x)
        model = keras.Model(inputs, outputs)

        optimizer = keras.optimizers.Adam()
        model.compile(loss="binary_crossentropy", optimizer=optimizer,
metrics=["accuracy"])
        return model, conv_base # Return conv_base for fine-tuning
access

```

```

except Exception as e:
    print(f"Error building pretrained Xception model: {e}")

    raise

def train_and_evaluate(model, train_dataset, val_dataset,
test_dataset, model_name, epochs=50):
    callbacks = [
        keras.callbacks.ModelCheckpoint(f"{model_name}.keras",
save_best_only=True, monitor="val_loss", verbose=0),
        keras.callbacks.EarlyStopping(monitor="val_loss", patience=10,
restore_best_weights=True, verbose=1) # Verbose early stopping
    ]
    print(f"Starting training for {model_name}...")
    history = model.fit(
        train_dataset,
        epochs=epochs,
        validation_data=val_dataset,
        callbacks=callbacks,
        verbose=1
    )
    print(f"Training finished for {model_name}. Loading best
model...")

    try:
        best_model = keras.models.load_model(f"{model_name}.keras")
        print("Best model loaded successfully.")
    except Exception as e:
        print(f"Could not load the best model from {model_name}.keras.
Using the model state from the end of training. Error: {e}")
        best_model = model

    print(f"Evaluating {model_name} on test data...")
    test_loss, test_acc = best_model.evaluate(test_dataset, verbose=0)
    print(f"Evaluation complete for {model_name}.")
    return {"history": history.history, "test_loss": test_loss,
"test_accuracy": test_acc}

def plot_training_history(history, title):
    if not history:
        print(f"No history data to plot for {title}")
        return
    acc = history.get("accuracy", [])
    val_acc = history.get("val_accuracy", [])
    loss = history.get("loss", [])
    val_loss = history.get("val_loss", [])

    if not acc or not val_acc or not loss or not val_loss:
        print(f"History data is incomplete for {title}. Cannot plot.")

```

```

        # print(f"Available keys: {history.keys()}") # Debugging line
        return

    epochs = range(1, len(acc) + 1)
    fig, (ax1, ax2) = plt.subplots(1, 2, figsize=(14, 5))
    ax1.plot(epochs, acc, 'r--', label='Training accuracy');
    ax1.plot(epochs, val_acc, 'b-', label='Validation accuracy')
    ax1.set_title(f'{title} - Accuracy'); ax1.legend()
    ax1.grid(True, alpha=0.3)
    ax2.plot(epochs, loss, 'r--', label='Training loss');
    ax2.plot(epochs, val_loss, 'b-', label='Validation loss')
    ax2.set_title(f'{title} - Loss'); ax2.legend()
    ax2.grid(True, alpha=0.3)
    plt.tight_layout()
    plt.show()

batch_size = 32
image_size = (180, 180)

val_set_size_per_class = 250
test_set_size_per_class = 250

experimental_val_dir =
pathlib.Path(f"experimental_validation_{val_set_size_per_class*2}")
experimental_test_dir =
pathlib.Path(f"experimental_test_{test_set_size_per_class*2}")

print("Creating smaller validation and test sets as required by
instructions...")
create_experimental_validation_set(source_val_dir,
experimental_val_dir, val_set_size_per_class)
create_experimental_test_set(source_test_dir, experimental_test_dir,
test_set_size_per_class)

print("\nLoading fixed smaller validation and test datasets...")
validation_dataset = image_dataset_from_directory(
    experimental_val_dir,
    image_size=image_size,
    batch_size=batch_size
).cache().prefetch(buffer_size=tf.data.AUTOTUNE)

test_dataset = image_dataset_from_directory(
    experimental_test_dir,
    image_size=image_size,
    batch_size=batch_size

```

```

).cache().prefetch(buffer_size=tf.data.AUTOTUNE)

print(f"\nFixed smaller validation set loaded:
{tf.data.experimental.cardinality(validation_dataset)*batch_size}
images")
print(f"Fixed smaller test set loaded:
{tf.data.experimental.cardinality(test_dataset)*batch_size} images")
print("Fixed smaller validation and test datasets ready.")

Creating smaller validation and test sets as required by
instructions...
Created experimental validation set with 250 images per class.
Created experimental test set with 250 images per class.

Loading fixed smaller validation and test datasets...
Found 500 files belonging to 2 classes.
Found 500 files belonging to 2 classes.

Fixed smaller validation set loaded: 512 images
Fixed smaller test set loaded: 512 images
Fixed smaller validation and test datasets ready.

scratch_results = []
experimental_train_sizes = [500, 750, 1000] # 1000, 1500, 2000 total
training images

for train_size in experimental_train_sizes:
    print(f"\n{'='*20} Starting Scratch Training with {train_size*2}
samples {'='*20}")
    print(f"Training from scratch with an experimental set of
{train_size} samples per class")

    experimental_train_dir =
pathlib.Path(f"experimental_train_{train_size}")

    create_experimental_training_set(source_train_dir,
experimental_train_dir, train_size) #

    train_dataset = image_dataset_from_directory(
        experimental_train_dir,
        image_size=image_size,
        batch_size=batch_size
    )

    model_scratch = build_scratch_model(input_shape=(180, 180, 3)) #
    model_scratch.summary()

```

```

    prefetched_train_dataset =
train_dataset.cache().prefetch(buffer_size=tf.data.AUTOTUNE)

    results_scratch = train_and_evaluate(
        model=model_scratch,
        train_dataset=prefetched_train_dataset,
        val_dataset=validation_dataset,
        test_dataset=test_dataset,
        model_name=f"scratch_train_{train_size}",
        epochs=100
    )

    plot_training_history(results_scratch["history"], f"From Scratch -
{train_size*2} Training Samples")
    scratch_results.append({
        "train_size": train_size * 2,
        "val_accuracy":
max(results_scratch["history"].get("val_accuracy", [0])),
        "test_accuracy": results_scratch["test_accuracy"]
    })
    print(f"\nFinal Results (Scratch, Train Size {train_size*2}): Test
Accuracy = {results_scratch['test_accuracy']:.4f}")

```

===== Starting Scratch Training with 1000 samples

=====

Training from scratch with an experimental set of 500 samples per class

Created experimental training set with 500 images per class.

Found 1000 files belonging to 2 classes.

Model: "functional_1"

Layer (type)	Output Shape	
Param #		
input_layer (InputLayer)	(None, 180, 180, 3)	
0		
data_augmentation (Sequential)	(None, 180, 180, 3)	
0		
rescaling (Rescaling)	(None, 180, 180, 3)	

0				
		conv2d (Conv2D)	(None, 178, 178, 32)	
896				
		max_pooling2d (MaxPooling2D)	(None, 89, 89, 32)	
0				
		conv2d_1 (Conv2D)	(None, 87, 87, 64)	
18,496				
		max_pooling2d_1 (MaxPooling2D)	(None, 43, 43, 64)	
0				
		conv2d_2 (Conv2D)	(None, 41, 41, 128)	
73,856				
		max_pooling2d_2 (MaxPooling2D)	(None, 20, 20, 128)	
0				
		conv2d_3 (Conv2D)	(None, 18, 18, 256)	
295,168				
		max_pooling2d_3 (MaxPooling2D)	(None, 9, 9, 256)	
0				
		conv2d_4 (Conv2D)	(None, 7, 7, 512)	
1,180,160				
		global_average_pooling2d	(None, 512)	
0		(GlobalAveragePooling2D)		
		dropout (Dropout)	(None, 512)	
0				

dense (Dense)	(None, 1)	
513		

Total params: 1,569,089 (5.99 MB)

Trainable params: 1,569,089 (5.99 MB)

Non-trainable params: 0 (0.00 B)

Starting training for scratch_train_500...

Epoch 1/100

32/32 ————— 45s 775ms/step - accuracy: 0.4850 - loss: 0.6981 - val_accuracy: 0.5000 - val_loss: 0.6919

Epoch 2/100

32/32 ————— 10s 316ms/step - accuracy: 0.5190 - loss: 0.6933 - val_accuracy: 0.5000 - val_loss: 0.6881

Epoch 3/100

32/32 ————— 12s 349ms/step - accuracy: 0.5520 - loss: 0.6914 - val_accuracy: 0.5020 - val_loss: 0.6814

Epoch 4/100

32/32 ————— 10s 324ms/step - accuracy: 0.5590 - loss: 0.6808 - val_accuracy: 0.5700 - val_loss: 0.6710

Epoch 5/100

32/32 ————— 9s 292ms/step - accuracy: 0.5720 - loss: 0.6650 - val_accuracy: 0.5620 - val_loss: 0.6776

Epoch 6/100

32/32 ————— 11s 324ms/step - accuracy: 0.5740 - loss: 0.6609 - val_accuracy: 0.5780 - val_loss: 0.6739

Epoch 7/100

32/32 ————— 10s 326ms/step - accuracy: 0.5890 - loss: 0.6533 - val_accuracy: 0.5880 - val_loss: 0.6594

Epoch 8/100

32/32 ————— 10s 323ms/step - accuracy: 0.5790 - loss: 0.6491 - val_accuracy: 0.6240 - val_loss: 0.6680

Epoch 9/100

32/32 ————— 10s 326ms/step - accuracy: 0.5910 - loss: 0.6492 - val_accuracy: 0.6020 - val_loss: 0.6745

Epoch 10/100

32/32 ————— 20s 324ms/step - accuracy: 0.6000 - loss: 0.6410 - val_accuracy: 0.6360 - val_loss: 0.6597

Epoch 11/100

32/32 ————— 21s 322ms/step - accuracy: 0.6320 - loss: 0.6258 - val_accuracy: 0.6220 - val_loss: 0.6460

Epoch 12/100

32/32 ————— 20s 327ms/step - accuracy: 0.6640 - loss: 0.6256 - val_accuracy: 0.5880 - val_loss: 0.6888

Epoch 13/100

32/32 ————— 21s 333ms/step - accuracy: 0.6190 - loss:


0.6347 - val_accuracy: 0.6500 - val_loss: 0.6390
Epoch 14/100
32/32 _____ 10s 327ms/step - accuracy: 0.6600 - loss:
0.6166 - val_accuracy: 0.6520 - val_loss: 0.6339
Epoch 15/100
32/32 _____ 11s 333ms/step - accuracy: 0.6400 - loss:
0.6224 - val_accuracy: 0.6680 - val_loss: 0.6225
Epoch 16/100
32/32 _____ 10s 308ms/step - accuracy: 0.6690 - loss:
0.6064 - val_accuracy: 0.6760 - val_loss: 0.6313
Epoch 17/100
32/32 _____ 10s 329ms/step - accuracy: 0.6570 - loss:
0.6031 - val_accuracy: 0.6760 - val_loss: 0.5959
Epoch 18/100
32/32 _____ 10s 324ms/step - accuracy: 0.6500 - loss:
0.6210 - val_accuracy: 0.6940 - val_loss: 0.6253
Epoch 19/100
32/32 _____ 10s 324ms/step - accuracy: 0.6850 - loss:
0.5983 - val_accuracy: 0.6700 - val_loss: 0.6255
Epoch 20/100
32/32 _____ 10s 327ms/step - accuracy: 0.6690 - loss:
0.5985 - val_accuracy: 0.6880 - val_loss: 0.5937
Epoch 21/100
32/32 _____ 10s 329ms/step - accuracy: 0.6900 - loss:
0.5824 - val_accuracy: 0.6920 - val_loss: 0.5818
Epoch 22/100
32/32 _____ 10s 302ms/step - accuracy: 0.7070 - loss:
0.5786 - val_accuracy: 0.7020 - val_loss: 0.5949
Epoch 23/100
32/32 _____ 10s 324ms/step - accuracy: 0.6780 - loss:
0.5791 - val_accuracy: 0.7100 - val_loss: 0.5759
Epoch 24/100
32/32 _____ 20s 327ms/step - accuracy: 0.6800 - loss:
0.5807 - val_accuracy: 0.7260 - val_loss: 0.5684
Epoch 25/100
32/32 _____ 10s 327ms/step - accuracy: 0.7150 - loss:
0.5772 - val_accuracy: 0.6680 - val_loss: 0.5806
Epoch 26/100
32/32 _____ 10s 326ms/step - accuracy: 0.6840 - loss:
0.5725 - val_accuracy: 0.6760 - val_loss: 0.5788
Epoch 27/100
32/32 _____ 10s 328ms/step - accuracy: 0.7050 - loss:
0.5611 - val_accuracy: 0.7060 - val_loss: 0.5667
Epoch 28/100
32/32 _____ 20s 316ms/step - accuracy: 0.7060 - loss:
0.5521 - val_accuracy: 0.6580 - val_loss: 0.5911
Epoch 29/100
32/32 _____ 10s 328ms/step - accuracy: 0.6870 - loss:
0.5781 - val_accuracy: 0.7120 - val_loss: 0.5691

```
Epoch 30/100
32/32 _____ 11s 330ms/step - accuracy: 0.7150 - loss:
0.5525 - val_accuracy: 0.7020 - val_loss: 0.5631
Epoch 31/100
32/32 _____ 10s 326ms/step - accuracy: 0.7080 - loss:
0.5551 - val_accuracy: 0.6700 - val_loss: 0.6269
Epoch 32/100
32/32 _____ 21s 330ms/step - accuracy: 0.7290 - loss:
0.5415 - val_accuracy: 0.6880 - val_loss: 0.5776
Epoch 33/100
32/32 _____ 10s 325ms/step - accuracy: 0.7420 - loss:
0.5231 - val_accuracy: 0.7200 - val_loss: 0.5594
Epoch 34/100
32/32 _____ 10s 319ms/step - accuracy: 0.7560 - loss:
0.5153 - val_accuracy: 0.7200 - val_loss: 0.5588
Epoch 35/100
32/32 _____ 11s 331ms/step - accuracy: 0.7500 - loss:
0.5042 - val_accuracy: 0.7400 - val_loss: 0.5411
Epoch 36/100
32/32 _____ 20s 327ms/step - accuracy: 0.7500 - loss:
0.4988 - val_accuracy: 0.7220 - val_loss: 0.5547
Epoch 37/100
32/32 _____ 21s 331ms/step - accuracy: 0.7580 - loss:
0.4822 - val_accuracy: 0.6680 - val_loss: 0.7725
Epoch 38/100
32/32 _____ 11s 334ms/step - accuracy: 0.7820 - loss:
0.4928 - val_accuracy: 0.7540 - val_loss: 0.5469
Epoch 39/100
32/32 _____ 10s 302ms/step - accuracy: 0.7710 - loss:
0.4560 - val_accuracy: 0.7320 - val_loss: 0.5926
Epoch 40/100
32/32 _____ 11s 322ms/step - accuracy: 0.7860 - loss:
0.4490 - val_accuracy: 0.7340 - val_loss: 0.6043
Epoch 41/100
32/32 _____ 10s 326ms/step - accuracy: 0.7770 - loss:
0.4605 - val_accuracy: 0.7200 - val_loss: 0.6321
Epoch 42/100
32/32 _____ 10s 324ms/step - accuracy: 0.7900 - loss:
0.4503 - val_accuracy: 0.7280 - val_loss: 0.5814
Epoch 43/100
32/32 _____ 10s 323ms/step - accuracy: 0.7870 - loss:
0.4559 - val_accuracy: 0.7360 - val_loss: 0.5635
Epoch 44/100
32/32 _____ 10s 321ms/step - accuracy: 0.8110 - loss:
0.4229 - val_accuracy: 0.7240 - val_loss: 0.5933
Epoch 45/100
32/32 _____ 11s 334ms/step - accuracy: 0.7980 - loss:
0.4293 - val_accuracy: 0.7300 - val_loss: 0.5808
Epoch 45: early stopping
```

```

Restoring model weights from the end of the best epoch: 35.
Training finished for scratch_train_500. Loading best model...
Best model loaded successfully.
Evaluating scratch_train_500 on test data...
Evaluation complete for scratch_train_500.

```



Final Results (Scratch, Train Size 1000): Test Accuracy = 0.7180

===== Starting Scratch Training with 1500 samples

=====

Training from scratch with an experimental set of 750 samples per class

Created experimental training set with 750 images per class.

Found 1500 files belonging to 2 classes.

Model: "functional_5"

Layer (type)	Output Shape
Param #	
input_layer_2 (InputLayer)	(None, 180, 180, 3)
data_augmentation (Sequential)	(None, 180, 180, 3)
rescaling_1 (Rescaling)	(None, 180, 180, 3)

conv2d_5 (Conv2D)	(None, 178, 178, 32)	
896		
max_pooling2d_4 (MaxPooling2D)	(None, 89, 89, 32)	
0		
conv2d_6 (Conv2D)	(None, 87, 87, 64)	
18,496		
max_pooling2d_5 (MaxPooling2D)	(None, 43, 43, 64)	
0		
conv2d_7 (Conv2D)	(None, 41, 41, 128)	
73,856		
max_pooling2d_6 (MaxPooling2D)	(None, 20, 20, 128)	
0		
conv2d_8 (Conv2D)	(None, 18, 18, 256)	
295,168		
max_pooling2d_7 (MaxPooling2D)	(None, 9, 9, 256)	
0		
conv2d_9 (Conv2D)	(None, 7, 7, 512)	
1,180,160		
global_average_pooling2d_1	(None, 512)	
0		
(GlobalAveragePooling2D)		
dropout_1 (Dropout)	(None, 512)	
0		
dense_1 (Dense)	(None, 1)	
513		

Total params: 1,569,089 (5.99 MB)

Trainable params: 1,569,089 (5.99 MB)

Non-trainable params: 0 (0.00 B)

Starting training for scratch_train_750...

Epoch 1/100

47/47 ————— 27s 583ms/step - accuracy: 0.4713 - loss: 0.6950 - val_accuracy: 0.5160 - val_loss: 0.6931

Epoch 2/100

47/47 ————— 15s 313ms/step - accuracy: 0.4780 - loss: 0.6934 - val_accuracy: 0.5000 - val_loss: 0.6931

Epoch 3/100

47/47 ————— 15s 323ms/step - accuracy: 0.5027 - loss: 0.6933 - val_accuracy: 0.5000 - val_loss: 0.6930

Epoch 4/100

47/47 ————— 16s 337ms/step - accuracy: 0.5147 - loss: 0.6931 - val_accuracy: 0.5000 - val_loss: 0.6930

Epoch 5/100

47/47 ————— 15s 324ms/step - accuracy: 0.4787 - loss: 0.6933 - val_accuracy: 0.5540 - val_loss: 0.6929

Epoch 6/100

47/47 ————— 16s 333ms/step - accuracy: 0.5080 - loss: 0.6925 - val_accuracy: 0.5060 - val_loss: 0.6858

Epoch 7/100

47/47 ————— 16s 343ms/step - accuracy: 0.5620 - loss: 0.6916 - val_accuracy: 0.5120 - val_loss: 0.6916

Epoch 8/100

47/47 ————— 15s 321ms/step - accuracy: 0.4953 - loss: 0.6931 - val_accuracy: 0.5000 - val_loss: 0.6925

Epoch 9/100

47/47 ————— 16s 332ms/step - accuracy: 0.4993 - loss: 0.6937 - val_accuracy: 0.5000 - val_loss: 0.6932

Epoch 10/100

47/47 ————— 15s 324ms/step - accuracy: 0.5007 - loss: 0.6932 - val_accuracy: 0.5000 - val_loss: 0.6926

Epoch 11/100

47/47 ————— 16s 335ms/step - accuracy: 0.5047 - loss: 0.6919 - val_accuracy: 0.5140 - val_loss: 0.6914

Epoch 12/100

47/47 ————— 15s 320ms/step - accuracy: 0.5273 - loss: 0.6883 - val_accuracy: 0.5200 - val_loss: 0.6861

Epoch 13/100

47/47 ————— 15s 320ms/step - accuracy: 0.5440 - loss: 0.6833 - val_accuracy: 0.5280 - val_loss: 0.6786

Epoch 14/100

47/47 ————— 15s 320ms/step - accuracy: 0.5640 - loss: 0.6698 - val_accuracy: 0.5680 - val_loss: 0.6584
Epoch 15/100
47/47 ————— 16s 335ms/step - accuracy: 0.5580 - loss: 0.6790 - val_accuracy: 0.5360 - val_loss: 0.6735
Epoch 16/100
47/47 ————— 15s 319ms/step - accuracy: 0.5580 - loss: 0.6749 - val_accuracy: 0.5340 - val_loss: 0.6770
Epoch 17/100
47/47 ————— 15s 319ms/step - accuracy: 0.5767 - loss: 0.6666 - val_accuracy: 0.5380 - val_loss: 0.6696
Epoch 18/100
47/47 ————— 15s 317ms/step - accuracy: 0.6333 - loss: 0.6519 - val_accuracy: 0.6200 - val_loss: 0.6460
Epoch 19/100
47/47 ————— 16s 335ms/step - accuracy: 0.6480 - loss: 0.6252 - val_accuracy: 0.6600 - val_loss: 0.6206
Epoch 20/100
47/47 ————— 15s 320ms/step - accuracy: 0.6480 - loss: 0.6267 - val_accuracy: 0.6920 - val_loss: 0.6149
Epoch 21/100
47/47 ————— 15s 318ms/step - accuracy: 0.6560 - loss: 0.6155 - val_accuracy: 0.7000 - val_loss: 0.6134
Epoch 22/100
47/47 ————— 15s 328ms/step - accuracy: 0.6660 - loss: 0.6180 - val_accuracy: 0.7140 - val_loss: 0.6122
Epoch 23/100
47/47 ————— 16s 344ms/step - accuracy: 0.6607 - loss: 0.6189 - val_accuracy: 0.6700 - val_loss: 0.6135
Epoch 24/100
47/47 ————— 15s 321ms/step - accuracy: 0.6700 - loss: 0.6061 - val_accuracy: 0.6960 - val_loss: 0.6034
Epoch 25/100
47/47 ————— 15s 314ms/step - accuracy: 0.6673 - loss: 0.5982 - val_accuracy: 0.7060 - val_loss: 0.6037
Epoch 26/100
47/47 ————— 15s 320ms/step - accuracy: 0.6833 - loss: 0.5927 - val_accuracy: 0.7120 - val_loss: 0.5897
Epoch 27/100
47/47 ————— 21s 325ms/step - accuracy: 0.6860 - loss: 0.5862 - val_accuracy: 0.7120 - val_loss: 0.5846
Epoch 28/100
47/47 ————— 15s 319ms/step - accuracy: 0.6800 - loss: 0.5933 - val_accuracy: 0.7060 - val_loss: 0.6102
Epoch 29/100
47/47 ————— 16s 336ms/step - accuracy: 0.6940 - loss: 0.5882 - val_accuracy: 0.6860 - val_loss: 0.6300
Epoch 30/100
47/47 ————— 15s 319ms/step - accuracy: 0.6940 - loss:

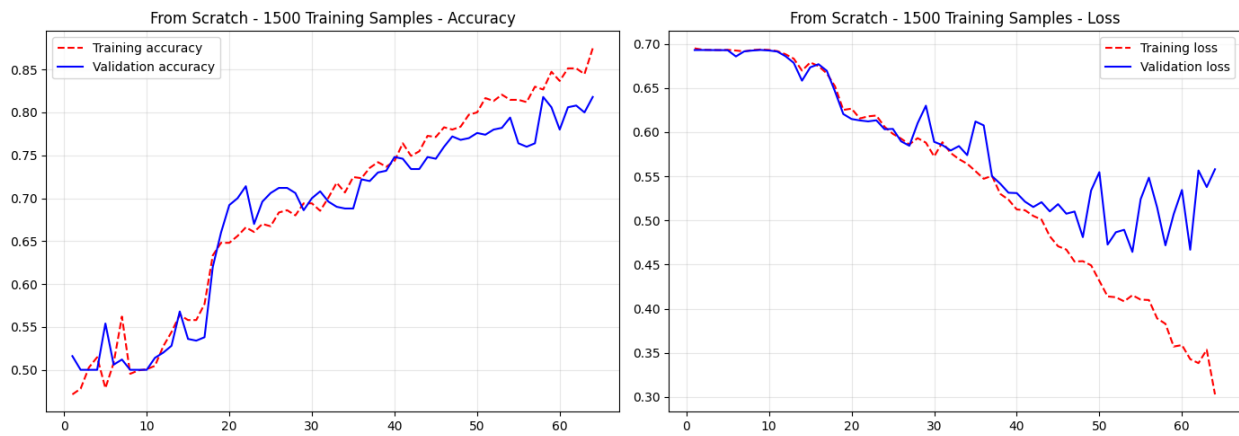
0.5725 - val_accuracy: 0.7000 - val_loss: 0.5891
Epoch 31/100
47/47 _____ 15s 315ms/step - accuracy: 0.6853 - loss: 0.5886 - val_accuracy: 0.7080 - val_loss: 0.5854
Epoch 32/100
47/47 _____ 21s 323ms/step - accuracy: 0.7007 - loss: 0.5765 - val_accuracy: 0.6960 - val_loss: 0.5793
Epoch 33/100
47/47 _____ 15s 316ms/step - accuracy: 0.7180 - loss: 0.5695 - val_accuracy: 0.6900 - val_loss: 0.5843
Epoch 34/100
47/47 _____ 15s 320ms/step - accuracy: 0.7067 - loss: 0.5643 - val_accuracy: 0.6880 - val_loss: 0.5740
Epoch 35/100
47/47 _____ 16s 332ms/step - accuracy: 0.7247 - loss: 0.5559 - val_accuracy: 0.6880 - val_loss: 0.6121
Epoch 36/100
47/47 _____ 15s 316ms/step - accuracy: 0.7233 - loss: 0.5472 - val_accuracy: 0.7220 - val_loss: 0.6077
Epoch 37/100
47/47 _____ 15s 319ms/step - accuracy: 0.7353 - loss: 0.5505 - val_accuracy: 0.7200 - val_loss: 0.5503
Epoch 38/100
47/47 _____ 15s 320ms/step - accuracy: 0.7420 - loss: 0.5302 - val_accuracy: 0.7300 - val_loss: 0.5417
Epoch 39/100
47/47 _____ 16s 342ms/step - accuracy: 0.7367 - loss: 0.5236 - val_accuracy: 0.7320 - val_loss: 0.5314
Epoch 40/100
47/47 _____ 20s 334ms/step - accuracy: 0.7433 - loss: 0.5125 - val_accuracy: 0.7480 - val_loss: 0.5310
Epoch 41/100
47/47 _____ 16s 342ms/step - accuracy: 0.7640 - loss: 0.5116 - val_accuracy: 0.7460 - val_loss: 0.5213
Epoch 42/100
47/47 _____ 16s 334ms/step - accuracy: 0.7493 - loss: 0.5052 - val_accuracy: 0.7340 - val_loss: 0.5152
Epoch 43/100
47/47 _____ 15s 319ms/step - accuracy: 0.7547 - loss: 0.5010 - val_accuracy: 0.7340 - val_loss: 0.5207
Epoch 44/100
47/47 _____ 15s 322ms/step - accuracy: 0.7727 - loss: 0.4821 - val_accuracy: 0.7480 - val_loss: 0.5103
Epoch 45/100
47/47 _____ 16s 334ms/step - accuracy: 0.7713 - loss: 0.4708 - val_accuracy: 0.7460 - val_loss: 0.5185
Epoch 46/100
47/47 _____ 20s 324ms/step - accuracy: 0.7827 - loss: 0.4670 - val_accuracy: 0.7600 - val_loss: 0.5077

Epoch 47/100
47/47 _____ 16s 346ms/step - accuracy: 0.7800 - loss: 0.4534 - val_accuracy: 0.7720 - val_loss: 0.5101
Epoch 48/100
47/47 _____ 15s 323ms/step - accuracy: 0.7833 - loss: 0.4539 - val_accuracy: 0.7680 - val_loss: 0.4811
Epoch 49/100
47/47 _____ 15s 319ms/step - accuracy: 0.7973 - loss: 0.4492 - val_accuracy: 0.7700 - val_loss: 0.5340
Epoch 50/100
47/47 _____ 15s 320ms/step - accuracy: 0.8000 - loss: 0.4315 - val_accuracy: 0.7760 - val_loss: 0.5547
Epoch 51/100
47/47 _____ 16s 339ms/step - accuracy: 0.8167 - loss: 0.4139 - val_accuracy: 0.7740 - val_loss: 0.4727
Epoch 52/100
47/47 _____ 15s 319ms/step - accuracy: 0.8133 - loss: 0.4130 - val_accuracy: 0.7800 - val_loss: 0.4867
Epoch 53/100
47/47 _____ 15s 320ms/step - accuracy: 0.8207 - loss: 0.4083 - val_accuracy: 0.7820 - val_loss: 0.4894
Epoch 54/100
47/47 _____ 15s 322ms/step - accuracy: 0.8147 - loss: 0.4153 - val_accuracy: 0.7940 - val_loss: 0.4644
Epoch 55/100
47/47 _____ 16s 337ms/step - accuracy: 0.8147 - loss: 0.4103 - val_accuracy: 0.7640 - val_loss: 0.5241
Epoch 56/100
47/47 _____ 15s 318ms/step - accuracy: 0.8120 - loss: 0.4099 - val_accuracy: 0.7600 - val_loss: 0.5485
Epoch 57/100
47/47 _____ 21s 338ms/step - accuracy: 0.8300 - loss: 0.3890 - val_accuracy: 0.7640 - val_loss: 0.5146
Epoch 58/100
47/47 _____ 15s 321ms/step - accuracy: 0.8267 - loss: 0.3829 - val_accuracy: 0.8180 - val_loss: 0.4718
Epoch 59/100
47/47 _____ 15s 317ms/step - accuracy: 0.8473 - loss: 0.3572 - val_accuracy: 0.8060 - val_loss: 0.5072
Epoch 60/100
47/47 _____ 21s 321ms/step - accuracy: 0.8367 - loss: 0.3588 - val_accuracy: 0.7800 - val_loss: 0.5345
Epoch 61/100
47/47 _____ 15s 322ms/step - accuracy: 0.8513 - loss: 0.3429 - val_accuracy: 0.8060 - val_loss: 0.4669
Epoch 62/100
47/47 _____ 15s 320ms/step - accuracy: 0.8513 - loss: 0.3384 - val_accuracy: 0.8080 - val_loss: 0.5566
Epoch 63/100

```

47/47 _____ 16s 343ms/step - accuracy: 0.8447 - loss:
0.3535 - val_accuracy: 0.8000 - val_loss: 0.5379
Epoch 64/100
47/47 _____ 15s 318ms/step - accuracy: 0.8747 - loss:
0.3029 - val_accuracy: 0.8180 - val_loss: 0.5580
Epoch 64: early stopping
Restoring model weights from the end of the best epoch: 54.
Training finished for scratch_train_750. Loading best model...
Best model loaded successfully.
Evaluating scratch_train_750 on test data...
Evaluation complete for scratch_train_750.

```



Final Results (Scratch, Train Size 1500): Test Accuracy = 0.7660

===== Starting Scratch Training with 2000 samples

Training from scratch with an experimental set of 1000 samples per class

Created experimental training set with 1000 images per class.

Found 2000 files belonging to 2 classes.

Model: "functional_9"

Layer (type) Param #	Output Shape
input_layer_3 (InputLayer) 0	(None, 180, 180, 3)
data_augmentation (Sequential) 0	(None, 180, 180, 3)

0	rescaling_2 (Rescaling)	(None, 180, 180, 3)
896	conv2d_10 (Conv2D)	(None, 178, 178, 32)
0	max_pooling2d_8 (MaxPooling2D)	(None, 89, 89, 32)
18,496	conv2d_11 (Conv2D)	(None, 87, 87, 64)
0	max_pooling2d_9 (MaxPooling2D)	(None, 43, 43, 64)
73,856	conv2d_12 (Conv2D)	(None, 41, 41, 128)
0	max_pooling2d_10 (MaxPooling2D)	(None, 20, 20, 128)
295,168	conv2d_13 (Conv2D)	(None, 18, 18, 256)
0	max_pooling2d_11 (MaxPooling2D)	(None, 9, 9, 256)
1,180,160	conv2d_14 (Conv2D)	(None, 7, 7, 512)
0	global_average_pooling2d_2 (GlobalAveragePooling2D)	(None, 512)
	dropout_2 (Dropout)	(None, 512)

0				
		dense_2 (Dense)	(None, 1)	
513				

Total params: 1,569,089 (5.99 MB)

Trainable params: 1,569,089 (5.99 MB)

Non-trainable params: 0 (0.00 B)

Starting training for scratch_train_1000...

Epoch 1/100

63/63 ————— 32s 515ms/step - accuracy: 0.5145 - loss: 0.6988 - val_accuracy: 0.5040 - val_loss: 0.6920

Epoch 2/100

63/63 ————— 20s 310ms/step - accuracy: 0.5365 - loss: 0.6903 - val_accuracy: 0.5560 - val_loss: 0.6841

Epoch 3/100

63/63 ————— 21s 321ms/step - accuracy: 0.5550 - loss: 0.6800 - val_accuracy: 0.6120 - val_loss: 0.6607

Epoch 4/100

63/63 ————— 20s 325ms/step - accuracy: 0.5875 - loss: 0.6678 - val_accuracy: 0.6320 - val_loss: 0.6379

Epoch 5/100

63/63 ————— 20s 312ms/step - accuracy: 0.6155 - loss: 0.6494 - val_accuracy: 0.6500 - val_loss: 0.6344

Epoch 6/100

63/63 ————— 20s 324ms/step - accuracy: 0.6170 - loss: 0.6532 - val_accuracy: 0.6540 - val_loss: 0.6390

Epoch 7/100

63/63 ————— 21s 330ms/step - accuracy: 0.6200 - loss: 0.6512 - val_accuracy: 0.6320 - val_loss: 0.6444

Epoch 8/100

63/63 ————— 20s 320ms/step - accuracy: 0.6300 - loss: 0.6448 - val_accuracy: 0.6380 - val_loss: 0.6418

Epoch 9/100

63/63 ————— 22s 335ms/step - accuracy: 0.6635 - loss: 0.6269 - val_accuracy: 0.6540 - val_loss: 0.6104

Epoch 10/100

63/63 ————— 21s 334ms/step - accuracy: 0.6660 - loss: 0.6173 - val_accuracy: 0.6580 - val_loss: 0.6311

Epoch 11/100

63/63 ————— 21s 336ms/step - accuracy: 0.6735 - loss: 0.6029 - val_accuracy: 0.6740 - val_loss: 0.6027

Epoch 12/100

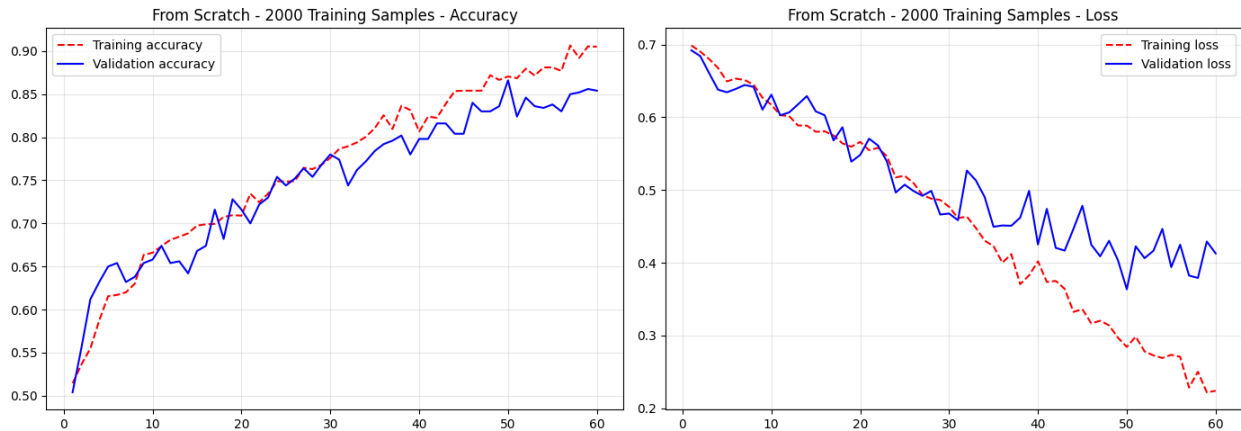
63/63 ————— 20s 313ms/step - accuracy: 0.6810 - loss:

0.6016 - val_accuracy: 0.6540 - val_loss: 0.6066
Epoch 13/100
63/63 ————— 20s 320ms/step - accuracy: 0.6845 - loss: 0.5887 - val_accuracy: 0.6560 - val_loss: 0.6177
Epoch 14/100
63/63 ————— 20s 322ms/step - accuracy: 0.6885 - loss: 0.5885 - val_accuracy: 0.6420 - val_loss: 0.6291
Epoch 15/100
63/63 ————— 20s 313ms/step - accuracy: 0.6975 - loss: 0.5801 - val_accuracy: 0.6680 - val_loss: 0.6081
Epoch 16/100
63/63 ————— 20s 323ms/step - accuracy: 0.6990 - loss: 0.5807 - val_accuracy: 0.6740 - val_loss: 0.6028
Epoch 17/100
63/63 ————— 20s 317ms/step - accuracy: 0.6995 - loss: 0.5750 - val_accuracy: 0.7160 - val_loss: 0.5682
Epoch 18/100
63/63 ————— 20s 318ms/step - accuracy: 0.7075 - loss: 0.5641 - val_accuracy: 0.6820 - val_loss: 0.5863
Epoch 19/100
63/63 ————— 20s 323ms/step - accuracy: 0.7095 - loss: 0.5595 - val_accuracy: 0.7280 - val_loss: 0.5390
Epoch 20/100
63/63 ————— 20s 313ms/step - accuracy: 0.7090 - loss: 0.5662 - val_accuracy: 0.7160 - val_loss: 0.5482
Epoch 21/100
63/63 ————— 20s 321ms/step - accuracy: 0.7345 - loss: 0.5547 - val_accuracy: 0.7000 - val_loss: 0.5705
Epoch 22/100
63/63 ————— 21s 328ms/step - accuracy: 0.7245 - loss: 0.5580 - val_accuracy: 0.7220 - val_loss: 0.5612
Epoch 23/100
63/63 ————— 20s 323ms/step - accuracy: 0.7345 - loss: 0.5465 - val_accuracy: 0.7300 - val_loss: 0.5393
Epoch 24/100
63/63 ————— 21s 324ms/step - accuracy: 0.7490 - loss: 0.5174 - val_accuracy: 0.7540 - val_loss: 0.4965
Epoch 25/100
63/63 ————— 21s 328ms/step - accuracy: 0.7485 - loss: 0.5195 - val_accuracy: 0.7440 - val_loss: 0.5073
Epoch 26/100
63/63 ————— 20s 318ms/step - accuracy: 0.7490 - loss: 0.5095 - val_accuracy: 0.7520 - val_loss: 0.4986
Epoch 27/100
63/63 ————— 20s 320ms/step - accuracy: 0.7645 - loss: 0.4932 - val_accuracy: 0.7640 - val_loss: 0.4920
Epoch 28/100
63/63 ————— 20s 326ms/step - accuracy: 0.7630 - loss: 0.4880 - val_accuracy: 0.7540 - val_loss: 0.4988

Epoch 29/100
63/63 ————— 40s 317ms/step - accuracy: 0.7680 - loss: 0.4863 - val_accuracy: 0.7680 - val_loss: 0.4663
Epoch 30/100
63/63 ————— 21s 327ms/step - accuracy: 0.7760 - loss: 0.4769 - val_accuracy: 0.7800 - val_loss: 0.4676
Epoch 31/100
63/63 ————— 20s 324ms/step - accuracy: 0.7865 - loss: 0.4613 - val_accuracy: 0.7740 - val_loss: 0.4584
Epoch 32/100
63/63 ————— 20s 312ms/step - accuracy: 0.7895 - loss: 0.4634 - val_accuracy: 0.7440 - val_loss: 0.5267
Epoch 33/100
63/63 ————— 20s 323ms/step - accuracy: 0.7940 - loss: 0.4480 - val_accuracy: 0.7620 - val_loss: 0.5134
Epoch 34/100
63/63 ————— 20s 322ms/step - accuracy: 0.8005 - loss: 0.4304 - val_accuracy: 0.7720 - val_loss: 0.4905
Epoch 35/100
63/63 ————— 20s 315ms/step - accuracy: 0.8105 - loss: 0.4228 - val_accuracy: 0.7840 - val_loss: 0.4495
Epoch 36/100
63/63 ————— 21s 322ms/step - accuracy: 0.8255 - loss: 0.4001 - val_accuracy: 0.7920 - val_loss: 0.4512
Epoch 37/100
63/63 ————— 21s 325ms/step - accuracy: 0.8095 - loss: 0.4117 - val_accuracy: 0.7960 - val_loss: 0.4508
Epoch 38/100
63/63 ————— 19s 309ms/step - accuracy: 0.8365 - loss: 0.3705 - val_accuracy: 0.8020 - val_loss: 0.4619
Epoch 39/100
63/63 ————— 20s 322ms/step - accuracy: 0.8315 - loss: 0.3825 - val_accuracy: 0.7800 - val_loss: 0.4988
Epoch 40/100
63/63 ————— 20s 323ms/step - accuracy: 0.8065 - loss: 0.4019 - val_accuracy: 0.7980 - val_loss: 0.4250
Epoch 41/100
63/63 ————— 20s 315ms/step - accuracy: 0.8240 - loss: 0.3736 - val_accuracy: 0.7980 - val_loss: 0.4740
Epoch 42/100
63/63 ————— 20s 324ms/step - accuracy: 0.8225 - loss: 0.3749 - val_accuracy: 0.8160 - val_loss: 0.4205
Epoch 43/100
63/63 ————— 20s 323ms/step - accuracy: 0.8390 - loss: 0.3641 - val_accuracy: 0.8160 - val_loss: 0.4166
Epoch 44/100
63/63 ————— 19s 309ms/step - accuracy: 0.8535 - loss: 0.3324 - val_accuracy: 0.8040 - val_loss: 0.4467
Epoch 45/100

```
63/63 _____ 20s 323ms/step - accuracy: 0.8540 - loss:
0.3361 - val_accuracy: 0.8040 - val_loss: 0.4783
Epoch 46/100
63/63 _____ 21s 324ms/step - accuracy: 0.8540 - loss:
0.3165 - val_accuracy: 0.8400 - val_loss: 0.4248
Epoch 47/100
63/63 _____ 20s 314ms/step - accuracy: 0.8540 - loss:
0.3202 - val_accuracy: 0.8300 - val_loss: 0.4088
Epoch 48/100
63/63 _____ 21s 320ms/step - accuracy: 0.8720 - loss:
0.3140 - val_accuracy: 0.8300 - val_loss: 0.4301
Epoch 49/100
63/63 _____ 20s 324ms/step - accuracy: 0.8665 - loss:
0.2969 - val_accuracy: 0.8360 - val_loss: 0.4035
Epoch 50/100
63/63 _____ 20s 312ms/step - accuracy: 0.8705 - loss:
0.2842 - val_accuracy: 0.8660 - val_loss: 0.3634
Epoch 51/100
63/63 _____ 21s 322ms/step - accuracy: 0.8685 - loss:
0.2983 - val_accuracy: 0.8240 - val_loss: 0.4226
Epoch 52/100
63/63 _____ 21s 323ms/step - accuracy: 0.8795 - loss:
0.2780 - val_accuracy: 0.8460 - val_loss: 0.4063
Epoch 53/100
63/63 _____ 19s 310ms/step - accuracy: 0.8715 - loss:
0.2726 - val_accuracy: 0.8360 - val_loss: 0.4165
Epoch 54/100
63/63 _____ 20s 319ms/step - accuracy: 0.8810 - loss:
0.2690 - val_accuracy: 0.8340 - val_loss: 0.4465
Epoch 55/100
63/63 _____ 21s 329ms/step - accuracy: 0.8810 - loss:
0.2733 - val_accuracy: 0.8380 - val_loss: 0.3940
Epoch 56/100
63/63 _____ 20s 317ms/step - accuracy: 0.8770 - loss:
0.2707 - val_accuracy: 0.8300 - val_loss: 0.4248
Epoch 57/100
63/63 _____ 21s 321ms/step - accuracy: 0.9065 - loss:
0.2283 - val_accuracy: 0.8500 - val_loss: 0.3823
Epoch 58/100
63/63 _____ 20s 317ms/step - accuracy: 0.8920 - loss:
0.2501 - val_accuracy: 0.8520 - val_loss: 0.3791
Epoch 59/100
63/63 _____ 20s 309ms/step - accuracy: 0.9055 - loss:
0.2218 - val_accuracy: 0.8560 - val_loss: 0.4292
Epoch 60/100
63/63 _____ 20s 319ms/step - accuracy: 0.9050 - loss:
0.2239 - val_accuracy: 0.8540 - val_loss: 0.4128
Epoch 60: early stopping
Restoring model weights from the end of the best epoch: 50.
```


Training finished for scratch_train_1000. Loading best model...
Best model loaded successfully.
Evaluating scratch_train_1000 on test data...
Evaluation complete for scratch_train_1000.



Final Results (Scratch, Train Size 2000): Test Accuracy = 0.8460

```
pretrained_results = []
```

```
for train_size in experimental_train_sizes:
    print(f"\n Starting Pretrained Xception with {train_size*2} samples")
    print(f"Using pretrained Xception model with an experimental set of {train_size} samples per class")

    experimental_train_dir =
    pathlib.Path(f"experimental_train_{train_size}")

    if not experimental_train_dir.exists():
        print(f"ERROR: Experimental training set for size {train_size} not found. Run Part 1 first.")
        continue

    train_dataset = image_dataset_from_directory(
        experimental_train_dir,
        image_size=image_size,
        batch_size=batch_size
    )

    print("\nBuilding model for Phase 1 (Frozen Base)...")

    model_frozen, conv_base_frozen =
```

```

build_pretrained_xception_model(trainable_base=False)
model_frozen.summary()

print("\nPhase 1: Training with frozen base...")

prefetched_train_dataset =
train_dataset.cache().prefetch(buffer_size=tf.data.AUTOTUNE)

results_frozen = train_and_evaluate(
    model=model_frozen,
    train_dataset=prefetched_train_dataset,
    val_dataset=validation_dataset,
    test_dataset=test_dataset,
    model_name=f"pretrained_frozen_xception_train_{train_size}",
    epochs=30
)
print(f"Phase 1 (Frozen Xception) Test Accuracy:
{results_frozen['test_accuracy']:.4f}")

print("\nPhase 2: Fine-tuning Xception...")

try:
    model_finetune = keras.models.load_model(
        f"pretrained_frozen_xception_train_{train_size}.keras"
    )
    print("Best frozen model loaded successfully for fine-
tuning.")
except Exception as e:
    print(f"Could not load the best frozen model from
{f'pretrained_frozen_xception_train_{train_size}'.keras}. Skipping
fine-tuning for this size. Error: {e}")
    continue

try:
    conv_base_finetune =
model_finetune.get_layer("xception_backbone")
    conv_base_finetune.trainable = True
    print(f"Xception backbone '{conv_base_finetune.name}'
unfrozen.")

    for layer in conv_base_finetune.layers:
        if isinstance(layer, layers.BatchNormalization):
            layer.trainable = False
    print("Batch normalization layers within the backbone frozen

```

```

for fine-tuning.")

except ValueError:
    print("Could not find the 'xception_backbone' layer. Check
model structure. Skipping fine-tuning.")
    continue

model_finetune.compile(
    loss="binary_crossentropy",
    optimizer=keras.optimizers.Adam(learning_rate=1e-5),
    metrics=["accuracy"]
)
print("Model recompiled for fine-tuning with low learning rate.")
model_finetune.summary()

# Fine-tune the model
results_finetuned = train_and_evaluate(
    model=model_finetune,
    train_dataset=prefetched_train_dataset,
    val_dataset=validation_dataset,
    test_dataset=test_dataset,

model_name=f"pretrained_finetuned_xception_train_{train_size}",
    epochs=30
)

plot_training_history(
    results_finetuned["history"],
    f"Pretrained Xception (Fine-tuned) - {train_size*2} Samples"
) #

pretrained_results.append({
    "train_size": train_size * 2,
    "val_accuracy":
max(results_finetuned["history"].get("val_accuracy", [0])),
    "test_accuracy": results_finetuned["test_accuracy"]
})
print(f"\nFinal Results (Fine-tuned Xception, Train Size
{train_size*2}): Test Accuracy =
{results_finetuned['test_accuracy']:.4f}")

```

Starting Pretrained Xception with 1000 samples
Using pretrained Xception model with an experimental set of 500
samples per class
Found 1000 files belonging to 2 classes.

Building model for Phase 1 (Frozen Base)...

Model: "functional_13"

Layer (type) Param #	Output Shape	
input_layer_5 (InputLayer) 0	(None, 180, 180, 3)	
data_augmentation (Sequential) 0	(None, 180, 180, 3)	
image_converter (ImageConverter) 0	(None, 180, 180, 3)	
xception_backbone 20,861,480 (XceptionBackbone)	(None, 6, 6, 2048)	
global_average_pooling2d_3 0 (GlobalAveragePooling2D)	(None, 2048)	
dense_3 (Dense) 524,544	(None, 256)	
dropout_3 (Dropout) 0	(None, 256)	
dense_4 (Dense) 257	(None, 1)	

Total params: 21,386,281 (81.58 MB)

Trainable params: 21,331,753 (81.37 MB)

Non-trainable params: 54,528 (213.00 KB)

Phase 1: Training with frozen base...

Starting training for pretrained_frozen_xception_train_500...

Epoch 1/30

32/32 _____ 187s 4s/step - accuracy: 0.8580 - loss: 0.3240 - val_accuracy: 0.9260 - val_loss: 0.9085

Epoch 2/30

32/32 _____ 82s 3s/step - accuracy: 0.9360 - loss: 0.1573 - val_accuracy: 0.7900 - val_loss: 2.3162

Epoch 3/30

32/32 _____ 80s 3s/step - accuracy: 0.9240 - loss: 0.2302 - val_accuracy: 0.7460 - val_loss: 3.8728

Epoch 4/30

32/32 _____ 90s 3s/step - accuracy: 0.9420 - loss: 0.1630 - val_accuracy: 0.9020 - val_loss: 0.7400

Epoch 5/30

32/32 _____ 83s 3s/step - accuracy: 0.9530 - loss: 0.1396 - val_accuracy: 0.9060 - val_loss: 0.5448

Epoch 6/30

32/32 _____ 82s 3s/step - accuracy: 0.9500 - loss: 0.1156 - val_accuracy: 0.8460 - val_loss: 1.4128

Epoch 7/30

32/32 _____ 81s 3s/step - accuracy: 0.9680 - loss: 0.1158 - val_accuracy: 0.9100 - val_loss: 0.5735

Epoch 8/30

32/32 _____ 88s 3s/step - accuracy: 0.9760 - loss: 0.0932 - val_accuracy: 0.9260 - val_loss: 0.3436

Epoch 9/30

32/32 _____ 82s 3s/step - accuracy: 0.9840 - loss: 0.0510 - val_accuracy: 0.9320 - val_loss: 0.3744

Epoch 10/30

32/32 _____ 81s 3s/step - accuracy: 0.9690 - loss: 0.0877 - val_accuracy: 0.7560 - val_loss: 2.4329

Epoch 11/30

32/32 _____ 82s 3s/step - accuracy: 0.9640 - loss: 0.1051 - val_accuracy: 0.9100 - val_loss: 0.3605

Epoch 12/30

32/32 _____ 81s 3s/step - accuracy: 0.9760 - loss: 0.0770 - val_accuracy: 0.8660 - val_loss: 0.5219

Epoch 13/30

32/32 _____ 82s 3s/step - accuracy: 0.9760 - loss: 0.0588 - val_accuracy: 0.9120 - val_loss: 0.3823

Epoch 14/30

32/32 _____ 89s 3s/step - accuracy: 0.9790 - loss: 0.0580 - val_accuracy: 0.9420 - val_loss: 0.2497

Epoch 15/30

32/32 _____ 81s 3s/step - accuracy: 0.9810 - loss: 0.0507 - val_accuracy: 0.7780 - val_loss: 1.5789

```
Epoch 16/30
32/32 _____ 83s 3s/step - accuracy: 0.9880 - loss:
0.0475 - val_accuracy: 0.9200 - val_loss: 0.5141
Epoch 17/30
32/32 _____ 81s 3s/step - accuracy: 0.9880 - loss:
0.0333 - val_accuracy: 0.8640 - val_loss: 0.6853
Epoch 18/30
32/32 _____ 82s 3s/step - accuracy: 0.9730 - loss:
0.0727 - val_accuracy: 0.6900 - val_loss: 1.5271
Epoch 19/30
32/32 _____ 90s 3s/step - accuracy: 0.9830 - loss:
0.0598 - val_accuracy: 0.9240 - val_loss: 0.2331
Epoch 20/30
32/32 _____ 93s 3s/step - accuracy: 0.9920 - loss:
0.0271 - val_accuracy: 0.9520 - val_loss: 0.1644
Epoch 21/30
32/32 _____ 131s 3s/step - accuracy: 0.9810 - loss:
0.0581 - val_accuracy: 0.9400 - val_loss: 0.2257
Epoch 22/30
32/32 _____ 141s 3s/step - accuracy: 0.9930 - loss:
0.0295 - val_accuracy: 0.9380 - val_loss: 0.2817
Epoch 23/30
32/32 _____ 83s 3s/step - accuracy: 0.9880 - loss:
0.0359 - val_accuracy: 0.8660 - val_loss: 0.9514
Epoch 24/30
32/32 _____ 141s 3s/step - accuracy: 0.9670 - loss:
0.0996 - val_accuracy: 0.8680 - val_loss: 0.5151
Epoch 25/30
32/32 _____ 95s 3s/step - accuracy: 0.9870 - loss:
0.0317 - val_accuracy: 0.9560 - val_loss: 0.1506
Epoch 26/30
32/32 _____ 81s 3s/step - accuracy: 0.9940 - loss:
0.0161 - val_accuracy: 0.9360 - val_loss: 0.2694
Epoch 27/30
32/32 _____ 81s 3s/step - accuracy: 0.9600 - loss:
0.1167 - val_accuracy: 0.7720 - val_loss: 0.9381
Epoch 28/30
32/32 _____ 82s 3s/step - accuracy: 0.9690 - loss:
0.0814 - val_accuracy: 0.9160 - val_loss: 0.2760
Epoch 29/30
32/32 _____ 81s 3s/step - accuracy: 0.9910 - loss:
0.0285 - val_accuracy: 0.9620 - val_loss: 0.1509
Epoch 30/30
32/32 _____ 80s 3s/step - accuracy: 0.9890 - loss:
0.0226 - val_accuracy: 0.9200 - val_loss: 0.2772
Restoring model weights from the end of the best epoch: 25.
Training finished for pretrained_frozen_xception_train_500. Loading
best model...
Best model loaded successfully.
```

Evaluating pretrained_frozen_xception_train_500 on test data...
Evaluation complete for pretrained_frozen_xception_train_500.
Phase 1 (Frozen Xception) Test Accuracy: 0.9200

Phase 2: Fine-tuning Xception...

Best frozen model loaded successfully for fine-tuning.

Xception backbone 'xception_backbone' unfrozen.

Batch normalization layers within the backbone frozen for fine-tuning.

Model recompiled for fine-tuning with low learning rate.

Model: "functional_13"

Layer (type) Param #	Output Shape	
input_layer_5 (InputLayer)	(None, 180, 180, 3)	
data_augmentation (Sequential)	(None, 180, 180, 3)	
image_converter (ImageConverter)	(None, 180, 180, 3)	
xception_backbone 20,861,480 (XceptionBackbone)	(None, 6, 6, 2048)	
global_average_pooling2d_3 (GlobalAveragePooling2D)	(None, 2048)	
dense_3 (Dense) 524,544	(None, 256)	
dropout_3 (Dropout)	(None, 256)	

dense_4 (Dense)	(None, 1)	
257		

Total params: 21,386,281 (81.58 MB)

Trainable params: 21,277,225 (81.17 MB)

Non-trainable params: 109,056 (426.00 KB)

Starting training for pretrained_finetuned_xception_train_500...

Epoch 1/30

32/32 ————— 58s 2s/step - accuracy: 0.9870 - loss: 0.0343 - val_accuracy: 0.9660 - val_loss: 0.1387

Epoch 2/30

32/32 ————— 89s 3s/step - accuracy: 0.9940 - loss: 0.0191 - val_accuracy: 0.9600 - val_loss: 0.1326

Epoch 3/30

32/32 ————— 49s 2s/step - accuracy: 0.9920 - loss: 0.0242 - val_accuracy: 0.9620 - val_loss: 0.1379

Epoch 4/30

32/32 ————— 48s 2s/step - accuracy: 0.9900 - loss: 0.0309 - val_accuracy: 0.9500 - val_loss: 0.1532

Epoch 5/30

32/32 ————— 54s 2s/step - accuracy: 0.9930 - loss: 0.0161 - val_accuracy: 0.9580 - val_loss: 0.1317

Epoch 6/30

32/32 ————— 84s 3s/step - accuracy: 0.9930 - loss: 0.0182 - val_accuracy: 0.9620 - val_loss: 0.1267

Epoch 7/30

32/32 ————— 48s 2s/step - accuracy: 0.9960 - loss: 0.0099 - val_accuracy: 0.9600 - val_loss: 0.1283

Epoch 8/30

32/32 ————— 61s 2s/step - accuracy: 0.9940 - loss: 0.0164 - val_accuracy: 0.9680 - val_loss: 0.1202

Epoch 9/30

32/32 ————— 77s 2s/step - accuracy: 0.9940 - loss: 0.0150 - val_accuracy: 0.9720 - val_loss: 0.1188

Epoch 10/30

32/32 ————— 60s 2s/step - accuracy: 0.9960 - loss: 0.0113 - val_accuracy: 0.9660 - val_loss: 0.1152

Epoch 11/30

32/32 ————— 47s 1s/step - accuracy: 0.9960 - loss: 0.0109 - val_accuracy: 0.9560 - val_loss: 0.1247

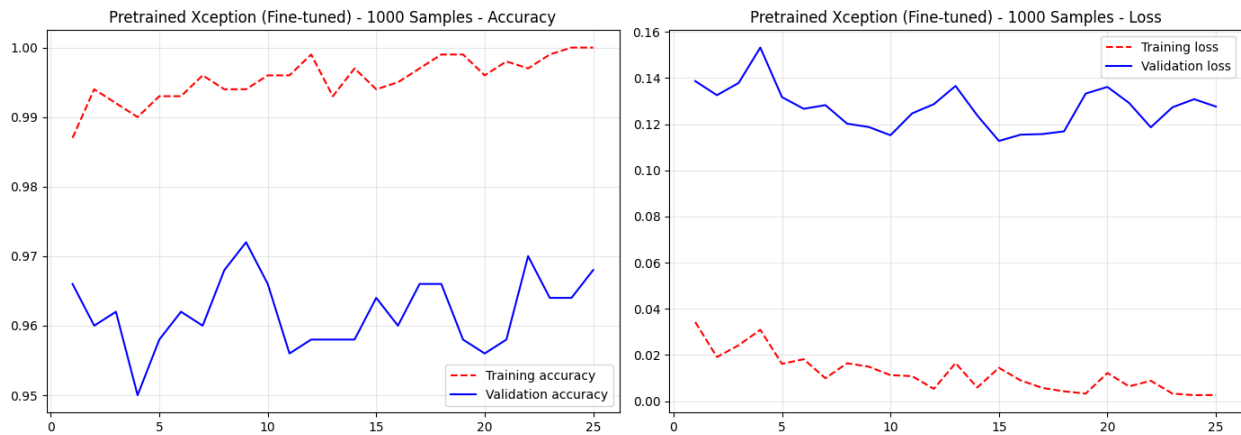
Epoch 12/30

32/32 ————— 48s 1s/step - accuracy: 0.9990 - loss: 0.0053 - val_accuracy: 0.9580 - val_loss: 0.1287


```
Epoch 13/30
32/32 _____ 48s 1s/step - accuracy: 0.9930 - loss:
0.0165 - val_accuracy: 0.9580 - val_loss: 0.1366
Epoch 14/30
32/32 _____ 48s 1s/step - accuracy: 0.9970 - loss:
0.0059 - val_accuracy: 0.9580 - val_loss: 0.1238
Epoch 15/30
32/32 _____ 79s 3s/step - accuracy: 0.9940 - loss:
0.0145 - val_accuracy: 0.9640 - val_loss: 0.1128
Epoch 16/30
32/32 _____ 50s 1s/step - accuracy: 0.9950 - loss:
0.0090 - val_accuracy: 0.9600 - val_loss: 0.1155
Epoch 17/30
32/32 _____ 47s 1s/step - accuracy: 0.9970 - loss:
0.0058 - val_accuracy: 0.9660 - val_loss: 0.1157
Epoch 18/30
32/32 _____ 48s 2s/step - accuracy: 0.9990 - loss:
0.0043 - val_accuracy: 0.9660 - val_loss: 0.1169
Epoch 19/30
32/32 _____ 81s 1s/step - accuracy: 0.9990 - loss:
0.0033 - val_accuracy: 0.9580 - val_loss: 0.1333
Epoch 20/30
32/32 _____ 82s 1s/step - accuracy: 0.9960 - loss:
0.0122 - val_accuracy: 0.9560 - val_loss: 0.1362
Epoch 21/30
32/32 _____ 48s 2s/step - accuracy: 0.9980 - loss:
0.0064 - val_accuracy: 0.9580 - val_loss: 0.1292
Epoch 22/30
32/32 _____ 48s 1s/step - accuracy: 0.9970 - loss:
0.0089 - val_accuracy: 0.9700 - val_loss: 0.1187
Epoch 23/30
32/32 _____ 48s 1s/step - accuracy: 0.9990 - loss:
0.0033 - val_accuracy: 0.9640 - val_loss: 0.1274
Epoch 24/30
32/32 _____ 83s 2s/step - accuracy: 1.0000 - loss:
0.0026 - val_accuracy: 0.9640 - val_loss: 0.1308
Epoch 25/30
32/32 _____ 47s 1s/step - accuracy: 1.0000 - loss:
0.0027 - val_accuracy: 0.9680 - val_loss: 0.1277
Epoch 25: early stopping
Restoring model weights from the end of the best epoch: 15.
Training finished for pretrained_finetuned_xception_train_500. Loading
best model...

/usr/local/lib/python3.12/dist-packages/keras/src/saving/
saving_lib.py:797: UserWarning: Skipping variable loading for
optimizer 'adam', because it has 318 variables whereas the saved
optimizer has 158 variables.
  saveable.load_own_variables(weights_store.get(inner_path))
```

Best model loaded successfully.
Evaluating pretrained_finetuned_xception_train_500 on test data...
Evaluation complete for pretrained_finetuned_xception_train_500.



Final Results (Fine-tuned Xception, Train Size 1000): Test Accuracy = 0.9500

Starting Pretrained Xception with 1500 samples
Using pretrained Xception model with an experimental set of 750 samples per class
Found 1500 files belonging to 2 classes.

Building model for Phase 1 (Frozen Base)...

Model: "functional_23"

Layer (type)	Output Shape
Param #	
input_layer_10 (InputLayer)	(None, 180, 180, 3)
data_augmentation (Sequential)	(None, 180, 180, 3)
image_converter	(None, 180, 180, 3)
(ImageConverter)	

20,861,480	xception_backbone	(None, 6, 6, 2048)	
	(XceptionBackbone)		
0	global_average_pooling2d_4	(None, 2048)	
	(GlobalAveragePooling2D)		
524,544	dense_5 (Dense)	(None, 256)	
0	dropout_4 (Dropout)	(None, 256)	
257	dense_6 (Dense)	(None, 1)	

Total params: 21,386,281 (81.58 MB)

Trainable params: 21,331,753 (81.37 MB)

Non-trainable params: 54,528 (213.00 KB)

Phase 1: Training with frozen base...

Starting training for pretrained_frozen_xception_train_750...

Epoch 1/30

47/47 ————— 173s 4s/step - accuracy: 0.8920 - loss: 0.2637 - val_accuracy: 0.8960 - val_loss: 1.2999

Epoch 2/30

47/47 ————— 155s 3s/step - accuracy: 0.9307 - loss: 0.1826 - val_accuracy: 0.9080 - val_loss: 0.7356

Epoch 3/30

47/47 ————— 127s 3s/step - accuracy: 0.9487 - loss: 0.1678 - val_accuracy: 0.9220 - val_loss: 0.4001

Epoch 4/30

47/47 ————— 120s 3s/step - accuracy: 0.9593 - loss: 0.1302 - val_accuracy: 0.9040 - val_loss: 0.7902

Epoch 5/30

47/47 ————— 130s 3s/step - accuracy: 0.9707 - loss: 0.0903 - val_accuracy: 0.9640 - val_loss: 0.1033

```

Epoch 6/30
47/47 _____ 118s 3s/step - accuracy: 0.9693 - loss:
0.0870 - val_accuracy: 0.9420 - val_loss: 0.5019
Epoch 7/30
47/47 _____ 143s 3s/step - accuracy: 0.9773 - loss:
0.0726 - val_accuracy: 0.9460 - val_loss: 0.2237
Epoch 8/30
47/47 _____ 142s 3s/step - accuracy: 0.9727 - loss:
0.0675 - val_accuracy: 0.8700 - val_loss: 0.5568
Epoch 9/30
47/47 _____ 121s 3s/step - accuracy: 0.9760 - loss:
0.0594 - val_accuracy: 0.9560 - val_loss: 0.1364
Epoch 10/30
47/47 _____ 119s 3s/step - accuracy: 0.9853 - loss:
0.0467 - val_accuracy: 0.9420 - val_loss: 0.1789
Epoch 11/30
47/47 _____ 119s 3s/step - accuracy: 0.9767 - loss:
0.0844 - val_accuracy: 0.8100 - val_loss: 1.1753
Epoch 12/30
47/47 _____ 118s 3s/step - accuracy: 0.9820 - loss:
0.0527 - val_accuracy: 0.8420 - val_loss: 0.5309
Epoch 13/30
47/47 _____ 142s 2s/step - accuracy: 0.9680 - loss:
0.1091 - val_accuracy: 0.5300 - val_loss: 6.5201
Epoch 14/30
47/47 _____ 118s 3s/step - accuracy: 0.9680 - loss:
0.0919 - val_accuracy: 0.9540 - val_loss: 0.2174
Epoch 15/30
47/47 _____ 118s 3s/step - accuracy: 0.9833 - loss:
0.0550 - val_accuracy: 0.9380 - val_loss: 0.3647
Epoch 15: early stopping
Restoring model weights from the end of the best epoch: 5.
Training finished for pretrained_frozen_xception_train_750. Loading
best model...
Best model loaded successfully.
Evaluating pretrained_frozen_xception_train_750 on test data...
Evaluation complete for pretrained_frozen_xception_train_750.
Phase 1 (Frozen Xception) Test Accuracy: 0.9720

```

```

Phase 2: Fine-tuning Xception...
Best frozen model loaded successfully for fine-tuning.
Xception backbone 'xception_backbone' unfrozen.
Batch normalization layers within the backbone frozen for fine-tuning.
Model recompiled for fine-tuning with low learning rate.

```

```

Model: "functional_23"

```

Layer (type)	Output Shape
--------------	--------------

Param #		
0	input_layer_10 (InputLayer)	(None, 180, 180, 3)
0	data_augmentation (Sequential)	(None, 180, 180, 3)
0	image_converter (ImageConverter)	(None, 180, 180, 3)
20,861,480	xception_backbone (XceptionBackbone)	(None, 6, 6, 2048)
0	global_average_pooling2d_4 (GlobalAveragePooling2D)	(None, 2048)
524,544	dense_5 (Dense)	(None, 256)
0	dropout_4 (Dropout)	(None, 256)
257	dense_6 (Dense)	(None, 1)

Total params: 21,386,281 (81.58 MB)

Trainable params: 21,277,225 (81.17 MB)

Non-trainable params: 109,056 (426.00 KB)

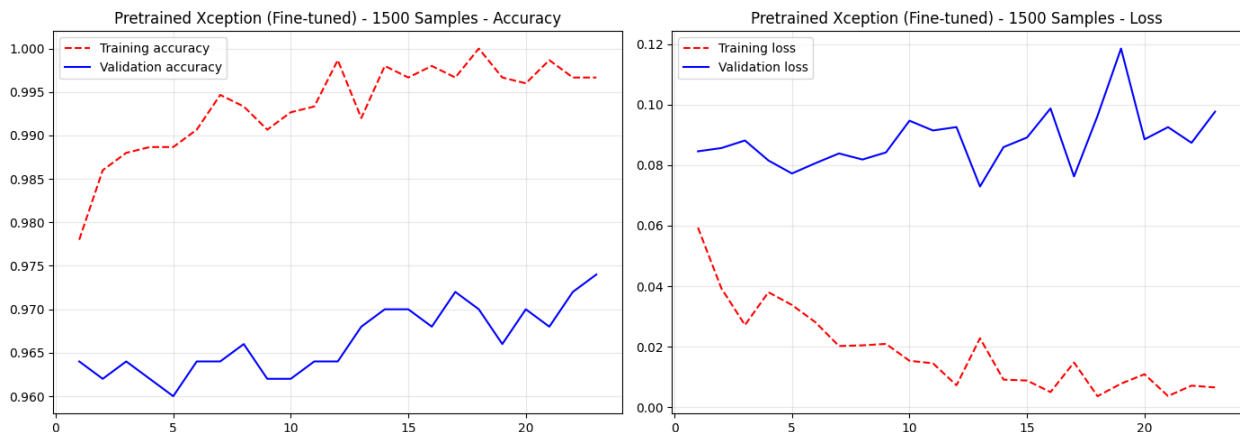
Starting training for pretrained_finetuned_xception_train_750...
Epoch 1/30

47/47 ————— 77s 2s/step - accuracy: 0.9780 - loss: 0.0594 - val_accuracy: 0.9640 - val_loss: 0.0846
Epoch 2/30
47/47 ————— 73s 1s/step - accuracy: 0.9860 - loss: 0.0393 - val_accuracy: 0.9620 - val_loss: 0.0856
Epoch 3/30
47/47 ————— 69s 1s/step - accuracy: 0.9880 - loss: 0.0272 - val_accuracy: 0.9640 - val_loss: 0.0881
Epoch 4/30
47/47 ————— 115s 2s/step - accuracy: 0.9887 - loss: 0.0380 - val_accuracy: 0.9620 - val_loss: 0.0815
Epoch 5/30
47/47 ————— 82s 2s/step - accuracy: 0.9887 - loss: 0.0338 - val_accuracy: 0.9600 - val_loss: 0.0772
Epoch 6/30
47/47 ————— 69s 1s/step - accuracy: 0.9907 - loss: 0.0281 - val_accuracy: 0.9640 - val_loss: 0.0806
Epoch 7/30
47/47 ————— 69s 1s/step - accuracy: 0.9947 - loss: 0.0202 - val_accuracy: 0.9640 - val_loss: 0.0839
Epoch 8/30
47/47 ————— 82s 1s/step - accuracy: 0.9933 - loss: 0.0205 - val_accuracy: 0.9660 - val_loss: 0.0818
Epoch 9/30
47/47 ————— 68s 1s/step - accuracy: 0.9907 - loss: 0.0209 - val_accuracy: 0.9620 - val_loss: 0.0842
Epoch 10/30
47/47 ————— 69s 1s/step - accuracy: 0.9927 - loss: 0.0154 - val_accuracy: 0.9620 - val_loss: 0.0947
Epoch 11/30
47/47 ————— 69s 1s/step - accuracy: 0.9933 - loss: 0.0145 - val_accuracy: 0.9640 - val_loss: 0.0914
Epoch 12/30
47/47 ————— 68s 1s/step - accuracy: 0.9987 - loss: 0.0073 - val_accuracy: 0.9640 - val_loss: 0.0926
Epoch 13/30
47/47 ————— 98s 2s/step - accuracy: 0.9920 - loss: 0.0228 - val_accuracy: 0.9680 - val_loss: 0.0729
Epoch 14/30
47/47 ————— 69s 1s/step - accuracy: 0.9980 - loss: 0.0091 - val_accuracy: 0.9700 - val_loss: 0.0859
Epoch 15/30
47/47 ————— 82s 1s/step - accuracy: 0.9967 - loss: 0.0088 - val_accuracy: 0.9700 - val_loss: 0.0891
Epoch 16/30
47/47 ————— 69s 1s/step - accuracy: 0.9980 - loss: 0.0050 - val_accuracy: 0.9680 - val_loss: 0.0987
Epoch 17/30
47/47 ————— 82s 1s/step - accuracy: 0.9967 - loss:

```

0.0148 - val_accuracy: 0.9720 - val_loss: 0.0763
Epoch 18/30
47/47 ━━━━━━━━━━━ 68s 1s/step - accuracy: 1.0000 - loss:
0.0037 - val_accuracy: 0.9700 - val_loss: 0.0963
Epoch 19/30
47/47 ━━━━━━━━━━━ 69s 1s/step - accuracy: 0.9967 - loss:
0.0078 - val_accuracy: 0.9660 - val_loss: 0.1185
Epoch 20/30
47/47 ━━━━━━━━━━━ 69s 1s/step - accuracy: 0.9960 - loss:
0.0109 - val_accuracy: 0.9700 - val_loss: 0.0885
Epoch 21/30
47/47 ━━━━━━━━━━━ 69s 1s/step - accuracy: 0.9987 - loss:
0.0037 - val_accuracy: 0.9680 - val_loss: 0.0926
Epoch 22/30
47/47 ━━━━━━━━━━━ 69s 1s/step - accuracy: 0.9967 - loss:
0.0071 - val_accuracy: 0.9720 - val_loss: 0.0874
Epoch 23/30
47/47 ━━━━━━━━━━━ 69s 1s/step - accuracy: 0.9967 - loss:
0.0066 - val_accuracy: 0.9740 - val_loss: 0.0977
Epoch 23: early stopping
Restoring model weights from the end of the best epoch: 13.
Training finished for pretrained_finetuned_xception_train_750. Loading
best model...
Best model loaded successfully.
Evaluating pretrained_finetuned_xception_train_750 on test data...
Evaluation complete for pretrained_finetuned_xception_train_750.

```



Final Results (Fine-tuned Xception, Train Size 1500): Test Accuracy = 0.9700

Starting Pretrained Xception with 2000 samples
Using pretrained Xception model with an experimental set of 1000 samples per class
Found 2000 files belonging to 2 classes.

Building model for Phase 1 (Frozen Base)...

Model: "functional_33"

Layer (type) Param #	Output Shape	
input_layer_15 (InputLayer) 0	(None, 180, 180, 3)	
data_augmentation (Sequential) 0	(None, 180, 180, 3)	
image_converter 0 (ImageConverter)	(None, 180, 180, 3)	
xception_backbone 20,861,480 (XceptionBackbone)	(None, 6, 6, 2048)	
global_average_pooling2d_5 0 (GlobalAveragePooling2D)	(None, 2048)	
dense_7 (Dense) 524,544	(None, 256)	
dropout_5 (Dropout) 0	(None, 256)	
dense_8 (Dense) 257	(None, 1)	

Total params: 21,386,281 (81.58 MB)

Trainable params: 21,331,753 (81.37 MB)

Non-trainable params: 54,528 (213.00 KB)

Phase 1: Training with frozen base...

Starting training for pretrained_frozen_xception_train_1000...

Epoch 1/30

63/63 ————— 205s 3s/step - accuracy: 0.8795 - loss: 0.2961 - val_accuracy: 0.6740 - val_loss: 4.8891

Epoch 2/30

63/63 ————— 171s 3s/step - accuracy: 0.9250 - loss: 0.2140 - val_accuracy: 0.8500 - val_loss: 1.7193

Epoch 3/30

63/63 ————— 223s 3s/step - accuracy: 0.9425 - loss: 0.1644 - val_accuracy: 0.9360 - val_loss: 0.2293

Epoch 4/30

63/63 ————— 160s 3s/step - accuracy: 0.9645 - loss: 0.0986 - val_accuracy: 0.8660 - val_loss: 0.8030

Epoch 5/30

63/63 ————— 207s 3s/step - accuracy: 0.9685 - loss: 0.0923 - val_accuracy: 0.9400 - val_loss: 0.2803

Epoch 6/30

63/63 ————— 159s 3s/step - accuracy: 0.9665 - loss: 0.0988 - val_accuracy: 0.9460 - val_loss: 0.2851

Epoch 7/30

63/63 ————— 164s 3s/step - accuracy: 0.9760 - loss: 0.0655 - val_accuracy: 0.9440 - val_loss: 0.3207

Epoch 8/30

63/63 ————— 195s 3s/step - accuracy: 0.9675 - loss: 0.0794 - val_accuracy: 0.9400 - val_loss: 0.2275

Epoch 9/30

63/63 ————— 178s 3s/step - accuracy: 0.9780 - loss: 0.0773 - val_accuracy: 0.9560 - val_loss: 0.1614

Epoch 10/30

63/63 ————— 177s 3s/step - accuracy: 0.9695 - loss: 0.0716 - val_accuracy: 0.9660 - val_loss: 0.1231

Epoch 11/30

63/63 ————— 166s 3s/step - accuracy: 0.9825 - loss: 0.0508 - val_accuracy: 0.9580 - val_loss: 0.1388

Epoch 12/30

63/63 ————— 201s 3s/step - accuracy: 0.9835 - loss: 0.0503 - val_accuracy: 0.9640 - val_loss: 0.1543

Epoch 13/30

63/63 ————— 166s 3s/step - accuracy: 0.9855 - loss: 0.0385 - val_accuracy: 0.9620 - val_loss: 0.2313

Epoch 14/30

63/63 ————— 165s 3s/step - accuracy: 0.9730 - loss:

0.0729 - val_accuracy: 0.8920 - val_loss: 0.6263
Epoch 15/30
63/63 ————— 236s 3s/step - accuracy: 0.9800 - loss: 0.0654 - val_accuracy: 0.9660 - val_loss: 0.0997
Epoch 16/30
63/63 ————— 164s 3s/step - accuracy: 0.9910 - loss: 0.0391 - val_accuracy: 0.9620 - val_loss: 0.1069
Epoch 17/30
63/63 ————— 163s 3s/step - accuracy: 0.9865 - loss: 0.0362 - val_accuracy: 0.9540 - val_loss: 0.1516
Epoch 18/30
63/63 ————— 163s 3s/step - accuracy: 0.9815 - loss: 0.0602 - val_accuracy: 0.9680 - val_loss: 0.1236
Epoch 19/30
63/63 ————— 203s 3s/step - accuracy: 0.9850 - loss: 0.0401 - val_accuracy: 0.9580 - val_loss: 0.1033
Epoch 20/30
63/63 ————— 164s 3s/step - accuracy: 0.9895 - loss: 0.0369 - val_accuracy: 0.9220 - val_loss: 0.3902
Epoch 21/30
63/63 ————— 161s 3s/step - accuracy: 0.9905 - loss: 0.0348 - val_accuracy: 0.9520 - val_loss: 0.1986
Epoch 22/30
63/63 ————— 159s 3s/step - accuracy: 0.9890 - loss: 0.0259 - val_accuracy: 0.9220 - val_loss: 0.3595
Epoch 23/30
63/63 ————— 161s 3s/step - accuracy: 0.9910 - loss: 0.0297 - val_accuracy: 0.8940 - val_loss: 0.5220
Epoch 24/30
63/63 ————— 161s 3s/step - accuracy: 0.9860 - loss: 0.0427 - val_accuracy: 0.8120 - val_loss: 1.2303
Epoch 25/30
63/63 ————— 160s 3s/step - accuracy: 0.9855 - loss: 0.0502 - val_accuracy: 0.9500 - val_loss: 0.1928
Epoch 25: early stopping
Restoring model weights from the end of the best epoch: 15.
Training finished for pretrained_frozen_xception_train_1000. Loading best model...
Best model loaded successfully.
Evaluating pretrained_frozen_xception_train_1000 on test data...
Evaluation complete for pretrained_frozen_xception_train_1000.
Phase 1 (Frozen Xception) Test Accuracy: 0.9600

Phase 2: Fine-tuning Xception...

Best frozen model loaded successfully for fine-tuning.

Xception backbone 'xception_backbone' unfrozen.

Batch normalization layers within the backbone frozen for fine-tuning.

Model recompiled for fine-tuning with low learning rate.

Model: "functional_33"

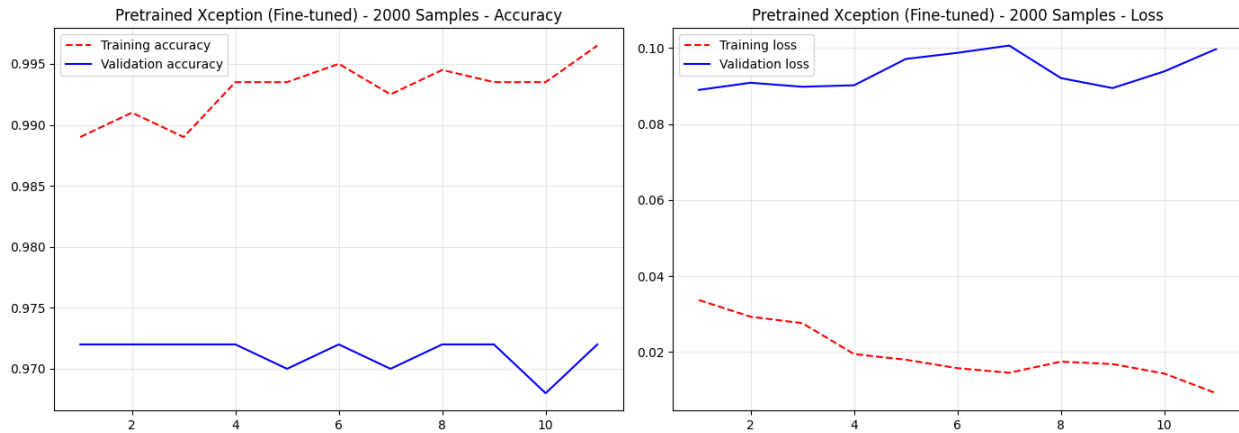
Layer (type) Param #	Output Shape	
input_layer_15 (InputLayer) 0	(None, 180, 180, 3)	
data_augmentation (Sequential) 0	(None, 180, 180, 3)	
image_converter 0 (ImageConverter)	(None, 180, 180, 3)	
xception_backbone 20,861,480 (XceptionBackbone)	(None, 6, 6, 2048)	
global_average_pooling2d_5 0 (GlobalAveragePooling2D)	(None, 2048)	
dense_7 (Dense) 524,544	(None, 256)	
dropout_5 (Dropout) 0	(None, 256)	
dense_8 (Dense) 257	(None, 1)	

Total params: 21,386,281 (81.58 MB)

Trainable params: 21,277,225 (81.17 MB)

Non-trainable params: 109,056 (426.00 KB)

```
Starting training for pretrained_finetuned_xception_train_1000...
Epoch 1/30
63/63 _____ 104s 2s/step - accuracy: 0.9890 - loss:
0.0337 - val_accuracy: 0.9720 - val_loss: 0.0890
Epoch 2/30
63/63 _____ 94s 2s/step - accuracy: 0.9910 - loss:
0.0293 - val_accuracy: 0.9720 - val_loss: 0.0909
Epoch 3/30
63/63 _____ 95s 2s/step - accuracy: 0.9890 - loss:
0.0276 - val_accuracy: 0.9720 - val_loss: 0.0898
Epoch 4/30
63/63 _____ 94s 1s/step - accuracy: 0.9935 - loss:
0.0195 - val_accuracy: 0.9720 - val_loss: 0.0902
Epoch 5/30
63/63 _____ 94s 1s/step - accuracy: 0.9935 - loss:
0.0180 - val_accuracy: 0.9700 - val_loss: 0.0971
Epoch 6/30
63/63 _____ 94s 1s/step - accuracy: 0.9950 - loss:
0.0157 - val_accuracy: 0.9720 - val_loss: 0.0987
Epoch 7/30
63/63 _____ 94s 1s/step - accuracy: 0.9925 - loss:
0.0145 - val_accuracy: 0.9700 - val_loss: 0.1006
Epoch 8/30
63/63 _____ 92s 1s/step - accuracy: 0.9945 - loss:
0.0174 - val_accuracy: 0.9720 - val_loss: 0.0921
Epoch 9/30
63/63 _____ 93s 1s/step - accuracy: 0.9935 - loss:
0.0168 - val_accuracy: 0.9720 - val_loss: 0.0895
Epoch 10/30
63/63 _____ 93s 1s/step - accuracy: 0.9935 - loss:
0.0143 - val_accuracy: 0.9680 - val_loss: 0.0939
Epoch 11/30
63/63 _____ 93s 1s/step - accuracy: 0.9965 - loss:
0.0092 - val_accuracy: 0.9720 - val_loss: 0.0997
Epoch 11: early stopping
Restoring model weights from the end of the best epoch: 1.
Training finished for pretrained_finetuned_xception_train_1000.
Loading best model...
Best model loaded successfully.
Evaluating pretrained_finetuned_xception_train_1000 on test data...
Evaluation complete for pretrained_finetuned_xception_train_1000.
```



Final Results (Fine-tuned Xception, Train Size 2000): Test Accuracy = 0.9660

```
if 'scratch_results' not in locals() or not scratch_results:
    print("\nScratch results not found. Please run Part 1.")
    df_scratch = pd.DataFrame(columns=['train_size', 'val_accuracy',
    'test_accuracy'])
else:
    df_scratch = pd.DataFrame(scratch_results)

if 'pretrained_results' not in locals() or not pretrained_results:
    print("Pretrained results not found. Please run Part 2.")
    df_pretrained = pd.DataFrame(columns=['train_size',
    'val_accuracy', 'test_accuracy'])
else:
    df_pretrained = pd.DataFrame(pretrained_results)

print("\n Scratch Results ")
print(df_scratch.to_string(index=False))

print("\nPretrained Model Results ")
print(df_pretrained.to_string(index=False))

fig, axes = plt.subplots(1, 2, figsize=(15, 6))

# Plot 1: Test Accuracy vs Training Size
axes[0].plot(df_scratch['train_size'], df_scratch['test_accuracy'],
            'o-', label='From Scratch', linewidth=2, markersize=8,
            color='red')
axes[0].plot(df_pretrained['train_size'],
            df_pretrained['test_accuracy'],
            's-', label='Pretrained (Xception)', linewidth=2,
            markersize=8, color='blue')
axes[0].set_xlabel('Training Set Size', fontsize=12,
```

```

fontweight='bold')
axes[0].set_ylabel('Test Accuracy', fontsize=12, fontweight='bold')
axes[0].set_title('Test Accuracy vs Training Set Size', fontsize=14,
fontweight='bold')
axes[0].legend(fontsize=11)
axes[0].grid(True, alpha=0.3)
axes[0].set_ylim([0.5, 1.0]) # Set y-axis from 50% to 100%

# Plot 2: Validation Accuracy vs Training Size
axes[1].plot(df_scratch['train_size'], df_scratch['val_accuracy'],
            'o-', label='From Scratch', linewidth=2, markersize=8,
            color='red')
axes[1].plot(df_pretrained['train_size'],
            df_pretrained['val_accuracy'],
            's-', label='Pretrained (Xception)', linewidth=2,
            markersize=8, color='blue')
axes[1].set_xlabel('Training Set Size', fontsize=12,
fontweight='bold')
axes[1].set_ylabel('Best Validation Accuracy', fontsize=12,
fontweight='bold')
axes[1].set_title('Validation Accuracy vs Training Set Size',
fontsize=14, fontweight='bold')
axes[1].legend(fontsize=11)
axes[1].grid(True, alpha=0.3)
axes[1].set_ylim([0.5, 1.0])

plt.tight_layout()
plt.savefig('comparison_results_xception_500valtest.png', dpi=300)
plt.show()

comparison_df = pd.merge(
    df_scratch[['train_size',
'test_accuracy']].rename(columns={'test_accuracy': 'Scratch Test
Acc'}),
    df_pretrained[['train_size',
'test_accuracy']].rename(columns={'test_accuracy': 'Pretrained Test
Acc'}),
    on='train_size',
    how='outer'
)
comparison_df['Improvement'] = (comparison_df['Pretrained Test Acc'] -
comparison_df['Scratch Test Acc'])

print("\n--- Comparison Table (Test Accuracy) ---")
print(comparison_df.round(4).to_string(index=False)) #

best_scratch_acc = 0
best_scratch_size = 'N/A'

```

```

if not df_scratch.empty:
    best_scratch_idx = df_scratch['test_accuracy'].idxmax()
    best_scratch_acc = df_scratch.loc[best_scratch_idx,
    'test_accuracy']
    best_scratch_size = df_scratch.loc[best_scratch_idx, 'train_size']

best_pretrained_acc = 0
best_pretrained_size = 'N/A'
if not df_pretrained.empty:
    best_pretrained_idx = df_pretrained['test_accuracy'].idxmax()
    best_pretrained_acc = df_pretrained.loc[best_pretrained_idx,
    'test_accuracy']
    best_pretrained_size = df_pretrained.loc[best_pretrained_idx,
    'train_size']

print("KEY FINDINGS / CONCLUSIONS (Based on this experiment)")

print(f"\nBest performance for training from scratch (among sizes
tested):")
print(f"  Training Size: {best_scratch_size}")
print(f"  Test Accuracy: {best_scratch_acc:.4f}")

print(f"\nBest performance for pretrained model (Xception, among sizes
tested):")
print(f"  Training Size: {best_pretrained_size}")
print(f"  Test Accuracy: {best_pretrained_acc:.4f}")



```

Scratch Results

train_size	val_accuracy	test_accuracy
1000	0.754	0.718
1500	0.818	0.766
2000	0.866	0.846

Pretrained Model Results

train_size	val_accuracy	test_accuracy
1000	0.972	0.950
1500	0.974	0.970
2000	0.972	0.966



--- Comparison Table (Test Accuracy) ---

train_size	Scratch Test Acc	Pretrained Test Acc	Improvement
1000	0.718	0.950	0.232
1500	0.766	0.970	0.204
2000	0.846	0.966	0.120

KEY FINDINGS / CONCLUSIONS (Based on this experiment)

Best performance for training from scratch (among sizes tested):

Training Size: 2000

Test Accuracy: 0.8460

Best performance for pretrained model (Xception, among sizes tested):

Training Size: 1500

Test Accuracy: 0.9700